

# Movie Rating Prediction - SIMPLE VERSION

## Easy-to-Understand Presentation

Student Name: [Your Name] Course: Machine Learning Date: November 2025

## ⊗ THE BIG IDEA (Read This First!)

### What Did We Build?

A system that predicts how much you'll like a movie BEFORE you watch it!

Think of it like this:

- Netflix shows you: "We think you'll rate this 4.5 stars"
- We built that! But simpler and for learning.

### Why Is This Important?

- Saves time - don't watch bad movies!
- Personalized - different people like different things
- Smart recommendations - like a friend who knows your taste

### What Makes This a PREDICTION Project?

- We predict NUMBERS (like 4.2 stars, 3.7 stars)
- NOT categories (not "good" vs "bad")
- Like predicting temperature (75°F) not weather (sunny/rainy)

## ⊗ SUPER SIMPLE OVERVIEW

### The 3 Main Things We Did:

#### 1⊗ Grouped Users by Behavior (K-Means)

Like sorting students into study groups:

- Group A: Love everything, watch tons of movies
- Group B: Picky, only watch a few movies
- Group C: Watch some, rate average
- Group D: Watch lots, but very picky

#### 2⊗ Built a Prediction Tree (Decision Tree)

Like a flowchart to predict ratings:

```
Is the movie popular?  
YES → Is the user generous with ratings?  
    YES → Predict 4.5 stars!  
    NO  → Predict 3.8 stars  
NO  → Is it rated highly by others?  
    YES → Predict 4.0 stars  
    NO  → Predict 2.5 stars
```

#### 3⊗ Found Hidden Patterns (Matrix Factorization)

Like finding out what people REALLY like:

- "This person likes action movies"
- "This movie is an action movie"
- Put them together → High rating!

# ¶ OUR DATA (The Movie Dataset)

## Quick Facts:

- 100,836 ratings from real people
- 610 users who rated movies
- 9,742 movies in the database
- Ratings scale: 0.5 to 5.0 stars (half star increments)

## The Big Problem:

98.3% of the data is MISSING!

Imagine a HUGE table:

- 610 rows (users)  $\times$  9,742 columns (movies) = 5,942,620 cells
- Only 100,836 have ratings (1.7%)
- Rest are empty because people haven't seen those movies

This is why we need smart algorithms!

# ¶ PHASE 1: Looking at the Data

## ¶ Quick Summary Box

What we did: Looked at the data to understand it Why: Can't build a model without knowing what we're working with Result: Found interesting patterns in ratings

## What We Discovered:

### Discovery 1: People Are Generous!

Most common rating: 4.0 stars (26.6% of all ratings!)  
Average rating: 3.50 stars  
People rarely give 0.5-2.0 stars (only 6% of ratings)

Why? People usually only rate movies they finished watching (and liked enough to finish!)

### Discovery 2: Some People Rate A LOT

Most active user: 2,698 ratings!  
Least active user: 20 ratings  
Average user: 165 ratings

Why it matters: Active users give us more data to learn from!

### Discovery 3: Most Movies Are Unknown

Popular movies: 300+ ratings (Forrest Gump, Shawshank Redemption)  
Most movies: Only 1-4 ratings (57% of all movies!)

The challenge: Hard to predict ratings for movies nobody's seen!

# ¶ PHASE 2: Grouping Users (K-Means)

## ¶ Quick Summary Box

What we did: Sorted users into 4 groups based on behavior Why: Different people rate differently How: Computer found patterns automatically Result: 4 clear groups of users

## The 4 User Groups We Found:

---

### Group 1: The Movie Lovers (25% of users)

- Watch 220+ movies
- Rate most things 3.7+ stars
- Think: Your friend who loves ALL movies

### Group 2: The Casual Viewers (30% of users)

- Watch only 45 movies
- Rate things 3.2 stars on average
- More critical/selective
- Think: Your friend who only watches blockbusters

### Group 3: The Regulars (28% of users)

- Watch 110 movies
- Rate things 3.5 stars (balanced)
- Think: Average movie watcher

### Group 4: The Critics (17% of users)

- Watch 180 movies
- Ratings are all over the place (1 to 5 stars)
- Think: The friend with strong opinions!

## How Did the Computer Find These Groups?

---

### Simple Explanation:

1. Measured 3 things for each user:
  - How many movies rated
  - Average rating they give
  - How much their ratings vary
2. Put each user as a point on a 3D graph
3. Computer found 4 clusters (groups) of nearby points
4. Users in same cluster behave similarly!

---

## ☒ PHASE 3: Building the Prediction Tree

### ☒ Quick Summary Box

**What we did:** Built a "decision tree" to predict ratings **Why:** Can make predictions for any user-movie combination **How:** Computer learned from 80,668 examples  
**Result:** Can predict ratings with ~0.9 star accuracy

### How the Decision Tree Works:

---

Think of it like 20 Questions for movie ratings!

Start here: What's the movie's average rating from others?

```
|— High ( $\geq 4.0$ )?  
| |— Is the user generous (avg rating  $\geq 4.0$ )?  
| | |— YES → Predict 4.5 stars! ⚡⚡⚡⚡  
| | |— NO → Predict 3.8 stars ⚡⚡⚡  
|  
|— Low ( $< 4.0$ )?  
| |— Is the movie popular (100+ ratings)?  
| | |— YES → Predict 3.5 stars ⚡⚡  
| | |— NO → Predict 2.5 stars ⚡
```

### Real Example:

---

**Question:** Will User #42 like "Die Hard"?

**Computer thinks:**

1. Die Hard has 4.2 average rating → High!
2. User #42 gives 4.3 stars on average → Generous!
3. **Prediction: 4.4 stars!** 🎉🎉

Actual rating User #42 gave: **4.5 stars** We were only 0.1 stars off! ✓

## How Accurate Is It?

---

**Our Results:**

- Average error: **0.71 stars**
- Root Mean Squared Error (RMSE): **0.92**

**What this means:**

- If we predict 4.0, actual rating is usually 3.3-4.7
  - Not perfect, but pretty good!
  - Much better than guessing!
- 

## ☒ PHASE 4: Finding Hidden Patterns (Matrix Factorization)

---

### ☒ Quick Summary Box

---

**What we did:** Found hidden "taste patterns" in the data **Why:** To fill in the 98.3% of missing ratings **How:** Math magic called SVD (Singular Value Decomposition)  
**Result:** Better predictions (0.88 RMSE)!

### The Big Idea (SUPER SIMPLE):

---

Imagine every person has invisible "taste scores":

- Action movie score: 0-100
- Comedy score: 0-100
- Drama score: 0-100
- Romance score: 0-100

And every movie has the same scores:

- How much action: 0-100
- How much comedy: 0-100
- How much drama: 0-100
- How much romance: 0-100

**To predict a rating:** Multiply person's scores × movie's scores = Predicted rating!

### Real Example (Simplified):

---

**User: Sarah**

- Loves action (90/100)
- Hates romance (10/100)
- Likes comedy (60/100)

**Movie: "Die Hard"**

- Lots of action (95/100)
- No romance (5/100)
- Some comedy (40/100)

**Prediction Math:**  $(90 \times 95) + (10 \times 5) + (60 \times 40) = \text{High score} \rightarrow \text{Predict 4.7 stars!}$

**Movie: "The Notebook"**

- No action (5/100)
- Lots of romance (95/100)
- No comedy (10/100)

**Prediction Math:**  $(90 \times 5) + (10 \times 95) + (60 \times 10) = \text{Low score} \rightarrow \text{Predict 2.1 stars!}$

### Why This Works Better:

---

**Decision Tree:**

- Uses only 7 features we created
- Can only split data in simple ways
- RMSE: 0.92

## Matrix Factorization:

- Finds 20 hidden patterns automatically
  - Captures complex relationships
  - RMSE: 0.88 (Better!)
- 

# ⊗ RESULTS: How Good Are We?

## ⊗ Quick Summary Box

Best Model: Matrix Factorization (SVD) Average Error: 0.68 stars Accuracy: Explains 35% of rating patterns

## The Report Card:

Model	Average Error	Grade
Matrix Factorization	0.68 stars	A-
Decision Tree	0.71 stars	B+
Random Guessing	1.25 stars	F

## What Does This Mean?

### ⊗ What We Can Do:

- Predict ratings usually within 0.7 stars
- Recommend movies you'll probably like
- Better than random guessing by 50%!

### ⊗ What We Can't Do (Yet):

- Perfect predictions (humans are unpredictable!)
- Predict for brand new users (no data yet)
- Predict for brand new movies (no ratings yet)
- Know what mood you're in today

## Why Only 35% Accuracy?

The other 65% is because:

- Personal mood (tired? stressed?)
- Who you're with (date vs friends vs alone)
- Random factors (someone talked during movie?)
- Individual quirks (some people just hate Brad Pitt!)

35% is actually pretty good for predicting human behavior!

---

# ⊗ WHAT DID WE LEARN?

## Key Finding #1: Movie Quality Matters Most

The movie's average rating from others is the #1 predictor (42% importance)

Translation: If everyone loves a movie, you'll probably like it too!

## Key Finding #2: User Groups Help

Grouping users by behavior improves predictions

Translation: People in the same group rate movies similarly!

## Key Finding #3: Hidden Patterns Are Powerful

Matrix Factorization beats simple decision trees

**Translation:** There are complex patterns humans can't see, but computers can!

## Key Finding #4: Data Sparsity Is Hard

---

98.3% missing data makes prediction challenging

**Translation:** We need LOTS of ratings to make good predictions!

---

## REAL WORLD USES

### Where Is This Used?

#### Netflix

"Because you watched..." **How:** Same algorithms, but MUCH bigger (billions of ratings!)

#### Spotify

"Discover Weekly" playlists **How:** Predict which songs you'll like based on patterns

#### Amazon

"Customers who bought this also bought..." **How:** Predict products you'll rate highly

#### YouTube

Recommended videos **How:** Predict which videos you'll watch and like

---

## COMMON QUESTIONS (Simple Answers!)

### Q: Why not just use movie genre?

A: Because people's tastes are complex! You might like action movies BUT only funny ones, OR only ones with a certain actor. Matrix factorization finds these hidden patterns automatically.

### Q: Why do we need 3 different methods?

A:

- **K-Means:** Groups users (unsupervised learning)
- **Decision Tree:** Makes predictions (supervised learning, interpretable)
- **Matrix Factorization:** Makes better predictions (supervised learning, accurate)

Each teaches us something different!

### Q: Can I use this for my own movie recommendations?

A: Yes! Just need to:

1. Rate 20+ movies yourself
2. Add your ratings to the dataset
3. Run the notebook
4. Get your personal recommendations!

### Q: How is this different from my friend's project (classification)?

A:

- **Their project:** "Is this email spam? YES or NO" (categories)
- **Our project:** "Rate this movie: 4.2 stars" (numbers)
- Different math, different evaluation, different algorithms!

### Q: What if I'm bad at math?

A: You don't need to understand ALL the math! Just the big ideas:

- K-Means = grouping similar things
- Decision Tree = if-then flowchart
- Matrix Factorization = finding hidden patterns

The computer does the hard math!

---

## ☒ THE BOTTOM LINE

### What We Built:

☒ Movie rating prediction system ☒ 3 different ML techniques ☒ Better than random guessing ☒ Real-world applicable

### What We Learned:

☒ How to group data (clustering) ☒ How to make predictions (regression) ☒ How to find patterns (matrix factorization) ☒ How to evaluate models (RMSE, MAE, R<sup>2</sup>)

### Why It Matters:

☒ This is how Netflix works! ☒ This is how Spotify works! ☒ This is how Amazon works! ☒ We built a mini version for learning!

---

## ☒ FINAL THOUGHTS

### The Journey:

1. Started with raw data (100K+ ratings)
2. Explored and understood the data
3. Created user groups
4. Built prediction models
5. Compared results
6. Found what works best!

### The Result:

We can predict movie ratings with ~0.7 star accuracy!

That means:

- If we predict 4.0 stars, actual is probably 3.3-4.7
  - Good enough to make useful recommendations
  - Shows we learned ML concepts successfully!
- 

## ☒ CHEAT SHEET: Key Numbers to Remember

Metric	Value	What It Means
Users	610	People who rated movies
Movies	9,742	Total movies in dataset
Ratings	100,836	Total ratings given
Sparsity	98.3%	How much data is missing
Clusters	4	User groups we found
Best RMSE	0.88	Average prediction error
Best R <sup>2</sup>	0.35	Variance explained

---

# ☰ PRESENTATION TIPS

---

## If Presenting:

---

### Start With:

"Have you ever wondered how Netflix knows what you'll like? We built that!"

### Main Points:

1. We predict NUMBERS not categories (regression not classification)
2. We used 3 techniques (clustering, trees, matrix factorization)
3. We achieved 0.7 star accuracy (pretty good!)
4. This is how real companies work (Netflix, Spotify, Amazon)

### End With:

"We turned math into movie magic! ☺"

---

# ☰ APPENDIX: Quick Definitions

---

**Clustering:** Grouping similar things together **Regression:** Predicting numbers (vs classification = predicting categories) **RMSE:** Root Mean Squared Error = average prediction error **MAE:** Mean Absolute Error = average distance from truth **R<sup>2</sup>:** How much variance we explain (0=bad, 1=perfect) **SVD:** Singular Value Decomposition = finding hidden patterns **Feature:** A measurable property (like "number of ratings") **Matrix:** A big table of numbers

---

### END OF SIMPLE PRESENTATION

*Remember: You don't need to understand every detail! Focus on the big ideas and you'll do great! ☺*

---

# ☰ For Your Presentation

---

### Use this version if:

- Your audience are beginners
- You want to keep it simple
- You want more analogies
- Time is limited (15-20 minutes)

### Use the detailed version if:

- Your audience knows ML
- Professor wants technical depth
- You need to show calculations
- Time is longer (30+ minutes)

**Best approach:** Use this simple version for slides, keep detailed version as backup for questions!